

# exp\_02

January 10, 2019

## 1 Data inspection

Below we can find the performance results of a mini benchmark `bench-stat` which were collected via `ci-hpc` framework.

The benchmark `bench-stat` was executed on a `charon` resource.

### 1.1 About `bench-stat` application

`bench-stat` application is a set of 3 benchmarks performing simple memory operations in a level 1, level 2 and level 3 cache.

These operations can be extremely fast, thus the experiments are *repeated*  $N$  times to obtain measurable duration of the benchmarks.

$$N = 1024 * 1024 * reps$$

where *reps* is a extra repetition coefficient which is altered in all of the commits.

As a baseline commit tagged as `reps-100` was selected, where *reps* = 100. The total number of repetition for this commit is  $N = 1024 * 1024 * 100 = 104\,857\,600$

A commits with maximum and minumum number of *reps* are tagged as `reps-125` and `reps-075` respectively.

### 1.2 Data structure

In table below we can see a *simplified* format of the data collected. The most of the fields are self-explanatory however some of them require explanation:

- `tag` - a git tag of a commit making the results more human-readable
- `timepoint` - numerical value of a tag for further purposes
- `no` - *i*-th repetition

```
In [1]: %matplotlib inline
        from cihpc.exp import exp_02_init as env
        env = env.reload(env)
        np, sc, pd, plt, sea = env.np, env.sc, env.pd, env.plt, env.sea
        df = env.fetch_data()
        df.head()
```

```
Out[1]:
```

	commit	walltime	time	tag	\$tag\$	no	walltime_mem_l1	\
0	75d38b25	5.420462	75	reps-075	\$reps^{075}\$	0	1.358015	
1	75d38b25	5.428641	75	reps-075	\$reps^{075}\$	1	1.373095	
2	75d38b25	5.359639	75	reps-075	\$reps^{075}\$	2	1.343519	
3	75d38b25	5.296926	75	reps-075	\$reps^{075}\$	3	1.343125	
4	75d38b25	5.333107	75	reps-075	\$reps^{075}\$	4	1.328733	

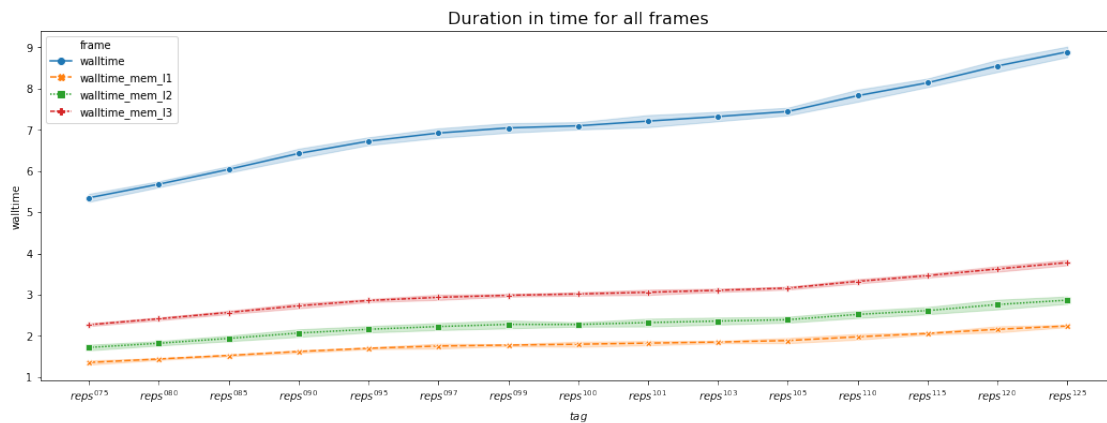
	walltime_mem_l2	walltime_mem_l3
0	1.808600	2.253832
1	1.766966	2.288569
2	1.784171	2.231938
3	1.682919	2.270869
4	1.763296	2.241055

### 1.3 Impact of individual commits on a duration

Chart below illustrates relation between walltime [sec] and commits, marked as

$reps^{075}, reps^{080}, \dots, reps^{095}, reps^{097}, reps^{099}, reps^{100}, reps^{101}, reps^{103}, reps^{105}, reps^{110}, \dots, reps^{125}$

```
In [2]: df2 = env.unwrap(df, ['walltime', 'walltime_mem_l1', 'walltime_mem_l2', 'walltime_mem_l3'])
plt.figure(figsize=(18, 6))
sea.lineplot(data=df2, x='$tag$', y='walltime', style='frame', hue='frame', markers=True)
plt.title('Duration in time for all frames', size=16);
```



### 1.4 Data distribution for individual commits

Charts below show histogram for each of the 15 commits along with normal fit (gray dashed line)

```
In [3]: g = sea.FacetGrid(df, col='tag', col_wrap=3, aspect=3, height=1.6, sharex=False, hue='tag')
g.map(sea.distplot, 'walltime', bins=11, rug=True, fit=sc.stats.norm, fit_kws=env.alpha)
g.fig.suptitle("Distribution of invidual commits", size=16)
g.fig.subplots_adjust(top=.9)
```

Distribution of individual commits

